

Claude Structured Outputs: Complete Implementation Guide

Overview

Structured Outputs guarantees that Claude's API responses match your exact JSON schema through constrained decoding. Released November 14, 2025 in public beta for Claude Sonnet 4.5 and Opus 4.1.

Key Benefit: Zero parsing errors, no retry logic, guaranteed schema compliance.

Quick Start

1. Beta Header (Required)

All requests must include:

```
headers = {
    "anthropic-beta": "structured-outputs-2025-11-13"
}
```

2. Two Modes Available

Mode	Use Case	Parameter
JSON Outputs	Data extraction, structured responses	<code>output_format</code>
Strict Tool Use	Type-safe function calls, agents	<code>strict: true</code>

Mode 1: JSON Outputs (Data Extraction)

Basic Usage

```
from anthropic import Anthropic

client = Anthropic()

schema = {
    "type": "object",
    "properties": {
        "name": {"type": "string"},
        "confidence": {"type": "number"},
        "category": {"type": "string", "enum": ["A", "B", "C"]}
    },
    "required": ["name", "confidence", "category"],
    "additionalProperties": False # Required for structured outputs
}

response = client.beta.messages.create(
    model="claude-sonnet-4-5-20250929",
    max_tokens=1024,
    betas=["structured-outputs-2025-11-13"],
    messages=[{
        "role": "user",
        "content": "Extract data from this text: ..."
    }],
    output_format={
        "type": "json_schema",
        "json_schema": {
            "name": "DataExtraction",
            "schema": schema,
            "strict": True
        }
    }
)

# Response is guaranteed valid JSON matching schema
data = response.content[0].text
```

SDK Helper: Pydantic Integration

Recommended approach - Let SDK handle schema transformation:

```
from anthropic import Anthropic
from pydantic import BaseModel

class ColorToken(BaseModel):
    hex: str
    confidence: float
    category: str

client = Anthropic()

response = client.beta.messages.parse(
    model="claude-sonnet-4-5-20250929",
    max_tokens=1024,
    betas=["structured-outputs-2025-11-13"],
    messages=[{"role": "user", "content": "Extract colors from this image"}],
    response_format=ColorToken
)

# Access validated data directly
color = response.parsed_output # Type: ColorToken
print(f"{color.hex} - {color.confidence}")
```

SDK Benefits:

- Automatic schema transformation (removes unsupported constraints)
- Built-in validation against original Pydantic model
- Returns typed `parsed_output` instead of raw JSON

Mode 2: Strict Tool Use (Agents & Function Calls)

Basic Usage

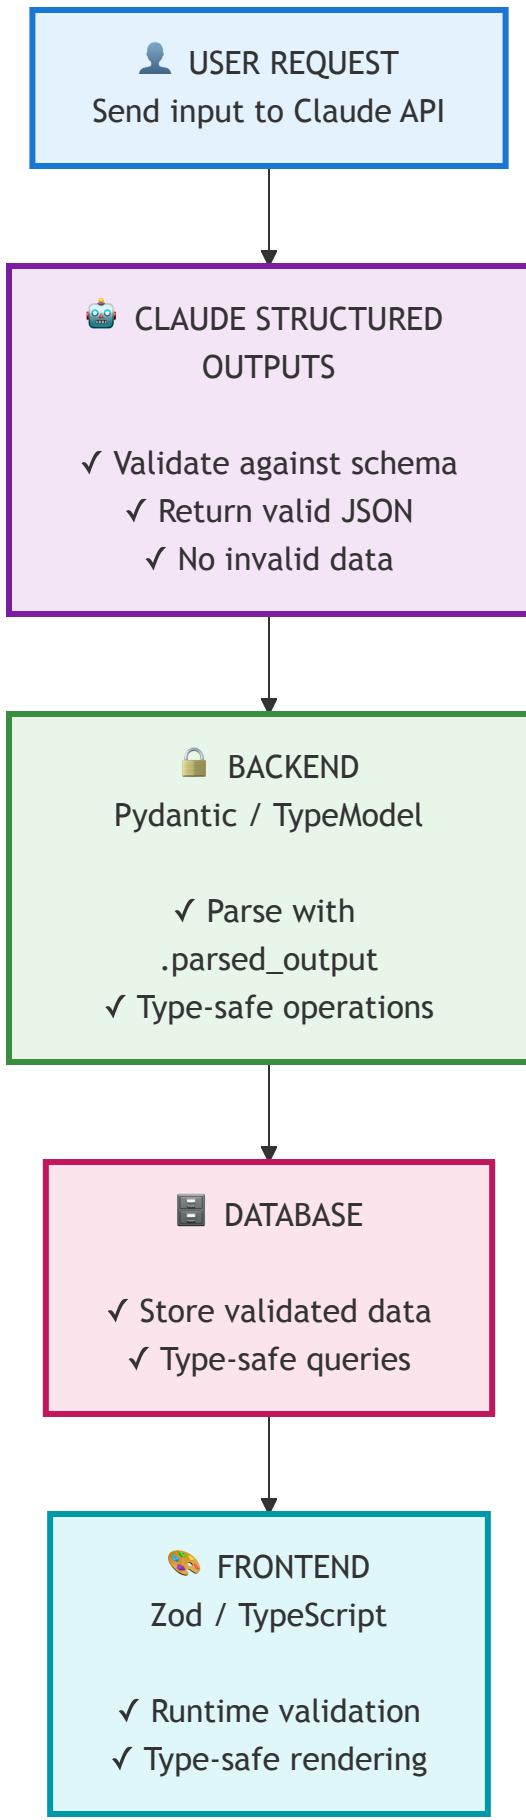
```
response = client.beta.messages.create(
    model="claude-sonnet-4-5-20250929",
    max_tokens=1024,
    betas=["structured-outputs-2025-11-13"],
    tools=[{
        "name": "get_weather",
        "description": "Get weather for a location",
        "input_schema": {
            "type": "object",
            "properties": {
                "location": {"type": "string"},
                "unit": {"type": "string", "enum": ["celsius", "fahrenheit"]}
            },
            "required": ["location", "unit"],
            "additionalProperties": False
        },
        "strict": True # Enables structured outputs for tools
    }],
    messages=[{"role": "user", "content": "What's the weather in Tokyo?"}]
)

# Tool input is guaranteed to match schema
if response.stop_reason == "tool_use":
    tool_call = response.content[0]
    # tool_call.input is type-safe and validated
```

Use Cases:

- Agentic workflows requiring reliable function calls
- API integrations where parameter validation is critical
- Multi-step reasoning with guaranteed tool compliance

Full-Stack Pipeline



TypeScript / Zod Integration

```
import Anthropic from '@anthropic-ai/sdk'
import { z } from 'zod'

const ColorSchema = z.object({
  hex: z.string(),
  confidence: z.number().min(0).max(1),
  category: z.enum(['primary', 'secondary', 'accent'])
})

const client = new Anthropic()

const response = await client.beta.messages.parse({
  model: 'claude-sonnet-4-5-20250929',
  max_tokens: 1024,
  betas: ['structured-outputs-2025-11-13'],
  messages: [{ role: 'user', content: 'Extract colors' }],
  response_format: ColorSchema
})

// Typed output
const color: z.infer<typeof ColorSchema> = response.parsed_output
```

Schema Requirements & Limitations

✓ Supported

- Basic types: string , number , boolean , integer , null , array , object
- Enums with string/number/boolean values
- Simple \$ref (within same schema)
- Nested objects and arrays

✗ Unsupported (SDK transforms automatically)

- Numerical constraints: minimum , maximum , multipleOf
- String constraints: minLength , maxLength , pattern
- Array constraints: minItems , maxItems , uniqueItems

- Recursive schemas or external \$ref
- additionalProperties: true (must be false)
- Complex enum types (e.g., object enums)

SDK Behavior: Unsupported constraints are removed during transformation, but validation occurs post-response using original schema.

Required Schema Properties

```
{
  "type": "object",
  "properties": { /* ... */ },
  "required": [ /* list required fields */ ],
  "additionalProperties": false // Must be false
}
```

Important Behaviors

1. Refusals Take Precedence

Safety/policy violations return stop_reason: "refusal" - response may not match schema:

```
if response.stop_reason == "refusal":
    print("Claude refused due to safety policies")
    # Handle gracefully, don't assume schema compliance
```

2. Token Limits

If max_tokens reached, output may be incomplete:

```
if response.stop_reason == "max_tokens":
    print("Increase max_tokens for complete response")
```

3. Grammar Caching (24 hours)

- First request with a schema: +100-200ms compilation latency

- Subsequent requests (24h): Near-zero overhead
- Compiled grammars cached per schema hash

4. Incompatibilities

Cannot use with:

- Citations (`stop_reason: "citation_error"`)
- Message prefilling
- Streaming responses (supported but no partial validation)

Cost Impact

- **Input tokens:** +0-3% (schema encoding overhead)
- **Output tokens:** Standard pricing
- **Batch API:** Works with 50% discount
- **Token counting:** Use `client.beta.messages.count_tokens()`

Savings: No retry logic or validation failures = net cost reduction

Best Practices

1. **Use SDK helpers** - `.parse()` simplifies everything
2. **Set additionalProperties: false** - Required for strict mode
3. **Include clear descriptions** - Helps Claude understand intent
4. **Provide examples in prompt** - "Return JSON like: {...}"
5. **Handle refusals gracefully** - Check `stop_reason`
6. **Monitor max_tokens** - Ensure sufficient budget for response
7. **Leverage grammar caching** - Reuse schemas for 24h performance boost

Copy That Implementation

Color Token Extraction (Real-World Example)

```
from anthropic import Anthropic
from pydantic import BaseModel, Field

class ColorToken(BaseModel):
    hex: str = Field(pattern=r'^#[0-9A-Fa-f]{6}$')
    confidence: float = Field(ge=0, le=1)
    semantic_name: str
    is_neutral: bool
    harmony: str

# AIColorExtractor using structured outputs
def extract_colors(image_data: bytes) -> list[ColorToken]:
    client = Anthropic()

    response = client.beta.messages.parse(
        model="claude-sonnet-4-5-20250929",
        max_tokens=2048,
        betas=["structured-outputs-2025-11-13"],
        messages=[{
            "role": "user",
            "content": [
                {
                    "type": "image",
                    "source": {
                        "type": "base64",
                        "media_type": "image/png",
                        "data": base64.b64encode(image_data).decode()
                    }
                },
                {
                    "type": "text",
                    "text": "Extract dominant colors with semantic names"
                }
            ]
        }],
        response_format=ColorToken
    )
```

```
# Guaranteed type-safe ColorToken
return response.parsed_output
```

Result: Zero Validation Overhead

- No manual hex validation
- No confidence bounds checking
- No field presence verification
- Direct database insertion
- Type-safe throughout stack

Architecture: Image → Claude (Structured Outputs) → Pydantic → Database → Zod → UI

Comparison: Before vs After

✗ Without Structured Outputs (100+ lines of defensive code)

```
def parse_claude_response(response_text: str) -> ColorToken:
    try:
        data = json.loads(response_text)
    except json.JSONDecodeError:
        raise ValidationError("Invalid JSON")

    if 'hex' not in data or not re.match(r'^#[0-9A-Fa-f]{6}$', data['hex']):
        raise ValidationError("Invalid hex")

    if 'confidence' not in data or not (0 <= data['confidence'] <= 1):
        raise ValidationError("Invalid confidence")

    # ... 50+ more lines of validation

    return ColorToken(**data)
```

✓ With Structured Outputs (1 line)

```
color = response.parsed_output # Guaranteed valid ColorToken
```

Additional Resources

- **Official Docs:** platform.claude.com/docs/en/build-with-claude/structured-outputs
- **Python SDK:** github.com/anthropics/anthropic-sdk-python
- **TypeScript SDK:** github.com/anthropics/anthropic-sdk-typescript

Summary

Structured Outputs = Zero parsing errors + Type safety + No retries

1. Add beta header: structured-outputs-2025-11-13
2. Use `.parse()` with Pydantic/Zod for best experience
3. Set `additionalProperties: false` in schemas
4. Handle refusals and token limits gracefully
5. Leverage 24-hour grammar caching for performance

Bottom Line: Define your schema once, get guaranteed compliance across your entire stack.